

AgentShield AI: An Autonomous Multi-Agent Framework for Syntactic Verification, Secret Interception, and Sandbox-Validated Remediation in Multi-Cloud Infrastructure-as-Code

Anisha Paturi

Dept. of Computer Science & Engg.
Keshav Memorial Inst. of Tech.
Hyderabad, Telangana, India
paturi.anisha@gmail.com

Ch Venkata Vahini

Dept. of Computer Science & Engg.
Keshav Memorial Inst. of Tech.
Hyderabad, Telangana, India
vahinivenkatac@gmail.com

Ch Parinamika Bhanu

Dept. of Computer Science & Engg.
Keshav Memorial Inst. of Tech.
Hyderabad, Telangana, India
chparinamikabhanu@gmail.com

Sravani Janak

Dept. of Computer Science & Engg.
Keshav Memorial Inst. of Tech.
Hyderabad, Telangana, India
sravanijanak@gmail.com

Abstract—Infrastructure-as-Code (IaC) has emerged as the foundational paradigm for declarative, automated, and scalable cloud resource provisioning across modern enterprise architectures. However, declarative configuration scripts—such as HashiCorp Terraform (HCL), AWS CloudFormation, Kubernetes Manifests, and Ansible Playbooks—introduce critical security vulnerabilities into continuous deployment pipelines. Common defects include over-privileged IAM policies, unencrypted data stores, unrestricted ingress security groups, and hardcoded cryptographic credentials. Existing static analysis security testing (SAST) tools rely on rigid regular expressions and superficial AST matching, yielding catastrophic false-positive rates (32%–48%) and demonstrating an inability to resolve cross-module variable interpolations or dynamic environment maps. Furthermore, conventional scanners are purely diagnostic, forcing human security engineers to manually author remediation patches, causing critical vulnerabilities to remain unresolved for an average Mean Time to Remediate (MTTR) of 24.6 days. In this paper, we present AgentShield AI, a fully autonomous, closed-loop multi-agent framework for zero-shot IaC security auditing, secret interception, and deterministic, sandbox-validated remediation. AgentShield AI deploys an ensemble of eight specialized autonomous agents coordinated by an event-driven Orchestration Router. The framework integrates Tree-sitter concrete syntax tree (CST) parsing, an entropy-aware dual secret interception engine combining Shannon entropy analysis (threshold $H \geq 4.5$) with 140+ heuristic regex patterns, and a hybrid Retrieval-Augmented Generation (RAG) architecture fusing Qdrant HNSW vector search over Center for Internet Security (CIS) Benchmarks with BM25 lexical ranking. Remediation patches are synthesized via dual-LLM consensus (Claude 3.5 Sonnet and GPT-4o) and strictly validated inside an isolated two-tier LocalStack Docker execution sandbox. Evaluated across 2,450 real-world IaC templates and the standardized Toprani-Madisetti benchmark, AgentShield AI achieves an unprecedented 99.1% detection precision, 98.4% recall, a 0.05% false-positive rate, a 97.8% sandbox-validated first-pass patch success rate, and an average execution latency of 1.84 seconds per module, significantly outperforming industry-standard baselines (Checkov, tfsec, KICS, Trivy).

Index Terms—Infrastructure-as-Code (IaC) Security, Multi-Agent Systems, Tree-sitter AST, Secret Detection, Shannon Entropy, Retrieval-Augmented Generation (RAG), LocalStack Sandbox, Automated Vulnerability Remediation, Cloud Compliance.

I. Introduction

Infrastructure-as-Code (IaC) has fundamentally transformed enterprise software engineering by enabling software-defined lifecycle management for distributed cloud architectures [1]. Through declarative domain-specific languages (DSLs) such as HashiCorp Terraform (HCL), AWS CloudFormation (JSON/YAML), Kubernetes Object Manifests, and Ansible Playbooks, engineering teams codify virtual networks, distributed databases, serverless execution fabrics, and identity access boundaries directly into version-controlled repositories [2], [3]. This programmatic representation allows continuous integration and continuous deployment (CI/CD) pipelines to provision, mutate, and tear down thousands of interconnected cloud assets in minutes, drastically reducing operational overhead, eliminating manual configuration drift, and standardizing infrastructure governance across heterogeneous cloud service providers [4].

However, this unprecedented provisioning velocity introduces profound security risks into enterprise software supply chains. Because IaC templates serve as executable blueprints for cloud perimeters, any configuration flaw—such as an Amazon S3 bucket missing public access blockades, an ingress firewall rule exposing port 22 (SSH) or 3389 (RDP) to the global Internet (0.0.0.0/0), an unencrypted elastic block store (EBS) volume, or an over-privileged IAM policy granting wildcard administrative capabilities—is immediately instantiated into live production infrastructure upon pipeline execution [5]. Empirical cloud threat intelligence reports indicate that over 73% of enterprise cloud security breaches originate from preventable IaC misconfigurations, while 65% of publicly accessible codebases contain hardcoded cryptographic credentials, secret tokens, or private RSA keys embedded directly within configuration variables [4].

Securing enterprise IaC pipelines presents four fundamental challenges that render existing commercial and academic tools inadequate:

1) High False Positive Rates in Static Scanners: Existing Static Application Security Testing (SAST) tools—such as Checkov [6], tfsec [7], KICS [8], and Trivy [9]—evaluate templates using rigid regular expressions or shallow Abstract Syntax Tree (AST) pattern matching. These scanners operate without semantic awareness of cross-file variable bindings, ternary conditional expressions, or hierarchical module inheritance, producing prohibitive false-alarm rates between 32% and 48% [10]. In enterprise environments managing

millions of lines of IaC, security engineers are overwhelmed by alert fatigue, frequently disabling automated guardrails to maintain release schedules.

2) Absence of Automated, Validated Remediation: Conventional SAST analyzers are strictly diagnostic; they output verbose violation logs without generating verified fix implementations [11]. Remediating these defects requires human cloud architects to inspect documentation, draft code patches, and test syntax manually, causing vulnerabilities to linger unpatched for an average Mean Time to Remediate (MTTR) of 24.6 days.

3) Unreliable and Hallucinatory LLM Generation: While recent generative AI models (such as GPT-4o and Claude 3.5 Sonnet) exhibit strong general coding abilities, direct zero-shot prompting for IaC repair suffers from severe hallucinations. Unconstrained LLMs frequently invent non-existent cloud resource attributes, violate strict provider schemas, deprecate valid tags, and introduce fatal syntax errors that break CI/CD execution [12], [13].

4) Cryptographic Secret Leakage: Hardcoded secrets require composite detection. Standard regex scanners miss novel token formats or obfuscated strings, while uncalibrated Shannon entropy analyzers generate massive false positives on benign UUIDs, commit hashes, and base64 assets [14].

To overcome these fundamental limitations, this paper presents **AgentShield AI**, an autonomous, closed-loop, multi-agent framework engineered for zero-shot IaC security verification, cryptographic secret interception, and deterministic, sandbox-validated remediation. AgentShield AI deploys eight specialized autonomous agents coordinated by an event-driven Orchestration Router. The framework couples Tree-sitter concrete syntax tree parsing with an entropy-aware dual secret interception engine, a hybrid dense-sparse Retrieval-Augmented Generation (RAG) system grounded in Center for Internet Security (CIS) Benchmarks [15], and an automated two-tier LocalStack Docker execution sandbox [16] that verifies operational semantics before emitting cryptographically signed Git Pull Requests.

II. Related Work

Securing Infrastructure-as-Code has evolved through several distinct paradigms, progressing from manual checklist audits to static rule engines, graph-theoretic dependency analyzers, formal SMT-based reasoning, and recent explorations in generative artificial intelligence.

A. Static Analysis and AST-Based Scanners: Checkov [6] builds intermediate Python AST representations to check Terraform and CloudFormation templates against CIS Benchmarks. tfsec [7] compiles HCL files into Go structures, evaluating static rules prior to plan generation. KICS [8] standardizes diverse IaC formats into a unified JSON model, evaluating rules via Open Policy Agent (OPA) Rego queries. Trivy [9] provides unified misconfiguration scanning across container images, Kubernetes manifests, and Terraform files. Despite their high throughput, these tools lack dynamic evaluation capabilities: they cannot resolve cross-module variable interpolation, module outputs, or ternary logical expressions, resulting in baseline false-positive rates between

32.4% and 47.9% [10]. Moreover, these tools provide no automated remediation mechanism.

B. Policy-as-Code and Formal SMT Verification: Policy-as-Code (PaC) frameworks, such as Open Policy Agent (OPA) and HashiCorp Sentinel, allow enterprise security teams to define declarative guardrails. However, authoring and maintaining complex Rego policies across evolving cloud APIs requires immense manual effort and domain expertise [10]. Formal verification approaches, including AWS Zelkova [17] and Cloud-SMR [18], translate access control policies into Satisfiability Modulo Theories (SMT) to mathematically prove the non-reachability of insecure states. While formal methods offer absolute theoretical soundness, they suffer from state explosion when applied to complex enterprise architectures containing thousands of interdependent resources, and they lack the generative capability to author corrective code.

C. Secret Interception and LLM Program Repair: Detecting exposed cryptographic keys in source code has traditionally relied on signature-based scanners such as TruffleHog [20] and Gitleaks [19]. Signature scanners match regular expressions for known key patterns (e.g., AWS AKIA prefixes, RSA private key headers). However, signature methods fail against custom tokens, base64-encoded secrets, or randomized strings. Shannon entropy analysis [14] measures the informational randomness of character strings, flagging tokens exhibiting high entropy. In the automated program repair domain, Toprani and Madiseti (2025) [21] introduced a graph-theoretic framework that leverages LLMs for Terraform security auditing. However, their architecture operates in an open-loop manner without closed-loop execution validation, leading to a 28.8% patch failure rate due to hallucinated attributes and broken provider dependencies.

III. System Architecture & Agent Methodology

AgentShield AI is architected as an event-driven, autonomous multi-agent ecosystem comprising eight specialized agents coordinated by a centralized Orchestration Router (Fig. 1). All agents execute asynchronously, exchanging strongly typed Pydantic V2 state objects across a shared execution bus:




Fig. 1. End-to-End System Architecture of AgentShield AI illustrating the 8-agent orchestration pipeline, Tree-sitter AST parsing, entropy-based secret scanning, hybrid RAG retrieval, dual-LLM consensus, and LocalStack sandbox validation.

• **Agent 1 (Orchestration & Ingestion Router):** Upon receiving an IaC repository or template snippet, Agent 1 identifies the source format (Terraform HCL, CloudFormation

JSON/YAML, Kubernetes Manifest, Ansible Playbook), extracts file hierarchy metadata, computes a SHA-256 integrity checksum, and initializes the shared execution context graph Gamma. The Router dynamically schedules parallel dispatch to Agent 2 and Agent 3 to minimize pipeline latency.

- **Agent 2 (AST & Graph-Theoretic Parser):** Converts raw declarative code into rich concrete syntax trees (CST) using high-performance C-bindings from Tree-sitter [22]. Agent 2 constructs an attributed directed graph $G = (V, E_{\text{dep}}, E_{\text{ref}}, A)$, resolving cross-block variable references, evaluating ternary conditional expressions, and identifying active resource nodes to eliminate false positives on disabled blocks.

- **Agent 3 (Dual-Engine Secret Interceptor):** Intercepts hardcoded API credentials, private certificates, and authentication tokens via a dual-stage analytical filter combining 140+ Gitleaks regex patterns with sliding-window Shannon entropy analysis [14] (threshold $H \geq 4.5$) and AST dictionary suppression.

- **Agent 4 (Hybrid RAG Knowledge Retrieval Engine):** Retrieves authoritative remediation guidelines from 12,400 chunked passages of CIS Cloud Benchmarks, NIST SP 800-53 Rev 5, and PCI-DSS v4.0 [15]. It fuses 384-dimensional dense semantic vectors (sentence-transformers/all-MiniLM-L6-v2) indexed in Qdrant HNSW graphs with sparse BM25 lexical ranking via Reciprocal Rank Fusion (RRF).

- **Agent 5 (Dual-LLM Consensus Remediation Generator):** Synthesizes candidate patches formatted as RFC 6902 JSON Patch arrays and Unified Git Diffs by dispatching structured prompts concurrently to Anthropic Claude 3.5 Sonnet (syntactic reasoning) and OpenAI GPT-4o (semantic synthesis). Patches are promoted only when token-level AST Dice consensus satisfies $S_{\text{dice}} \geq 0.92$.

- **Agent 6 (Two-Tier LocalStack Docker Sandbox Validator):** Validates patches inside an ephemeral, isolated Docker container running LocalStack (simulating 45+ AWS APIs) and local Terraform runtimes [16]. Tier 1 verifies AST syntax invariants (terraform validate); Tier 2 executes mock cloud provisioning (terraform plan/apply).

- **Agent 7 (Compliance Mapping & Threat Model Analyzer):** Enriches findings with Common Weakness Enumeration (CWE), CVSS v3.1 vector strings, CIS sub-controls, and MITRE ATT&CK; for Cloud matrix techniques (e.g., T1078 Valid Accounts, T1530 Data from Cloud Storage Object).

- **Agent 8 (Cryptographic Report & Git Pull Request Generator):** Aggregates execution artifacts into executive summaries, SARIF JSON files for GitHub/GitLab Security tab integration, and cryptographically signs Git Pull Requests using ephemeral Ed25519 keys.

IV. Mathematical Formulation & Algorithmic Workflow

To formalize the mathematical foundation of AgentShield AI, we define the core formulations for entropy detection, context fusion, consensus, and validation:

Shannon Entropy of candidate token string S of length L with character frequencies $f(c)$:

$$H(S) = - \sum_{i=1}^n P(c_i) \log_2 P(c_i) = - \sum_{i=1}^n \frac{n(c_i)}{L} \log_2 \left(\frac{n(c_i)}{L} \right) \quad (1)$$

Reciprocal Rank Fusion (RRF) score for document d across dense (Qdrant HNSW) and sparse (BM25) retrieval models (constant $k = 60$):

$$\text{RRF_Score}(d) = \sum_{m \in \{\text{Dense}, \text{Sparse}\}} \frac{1}{k + r_m(d)} \quad (2)$$

Token-level AST Dice similarity coefficient between candidate patches δ_1 (Claude 3.5) and δ_2 (GPT-4o):

$$S_{\text{dice}}(\delta_1, \delta_2) = \frac{2|\text{AST}(\delta_1) \cap \text{AST}(\delta_2)|}{|\text{AST}(\delta_1)| + |\text{AST}(\delta_2)|} \quad (3)$$

Two-tier LocalStack sandbox validation scoring function:

$$V_{\text{score}}(\delta) = 0.2 \cdot S_{\text{syntax}}(\delta) + 0.3 \cdot S_{\text{plan}}(\delta) + 0.5 \cdot S_{\text{apply}}(\delta) \quad (4)$$

where S_{syntax} , S_{plan} , and S_{apply} in $\{0, 1\}$ represent boolean execution success flags. Algorithm 1 formalizes the end-to-end multi-agent execution pipeline.

Algorithm 1: Autonomous Multi-Agent IaC Auditing and Sandbox Remediation
Input: IaC File T_{raw} ; Compliance Policy P_{cis}
Output: Signed SARIF Report R_{sarif} ; Validated Patch Δ_{final}
1: Initialize Shared Context $\Gamma \leftarrow \emptyset$; Compute Hash $H_0 \leftarrow \text{SHA256}(T_{\text{raw}})$;
2: [Agent 1] Detect Format; [Agent 2] Parse AST $G_{\text{AST}} \leftarrow \text{TreeSitterParse}(T_{\text{raw}})$;
3: [Agent 3] Compute Entropy $H(S_i)$; Intercept & redact secrets if $H(S_i) \geq 4.5$;
4: [Agent 2] Evaluate CIS Policy Rules; Identify Violation Set $V = \{v_1, \dots, v_K\}$;
5: **for each** violation $v_k \in V$ **do**
6: [Agent 4] Retrieve Context $C_k \leftarrow \text{RRF}(\text{QdrantHNSW}(v_k), \text{BM25}(v_k))$;
7: [Agent 5] Concurrently prompt Claude 3.5 Sonnet & GPT-4o; Compute S_{dice} ;
8: [Agent 6] Tier 1: AST Syntax Validation $S_{\text{syntax}}(\delta_k)$;
9: [Agent 6] Tier 2: LocalStack Docker Mock Provisioning $S_{\text{apply}}(\delta_k)$;
10: **if** $V_{\text{score}}(\delta_k) = 1.0$ **then** Accept Patch $\Delta_{\text{final}} \leftarrow \Delta_{\text{final}} \cup \delta_k$;
11: **else** Feed compiler error to Agent 5 for retry (max 3 iters);
12: [Agent 7] Map CWE/CVSS/CIS/ATT&CK; [Agent 8] Generate SARIF & Signed Git PR;
13: **return** $R_{\text{sarif}}, \Delta_{\text{final}}$

V. Experimental Setup & Benchmark Methodology

A. Benchmark Datasets: Evaluated across three suites totaling 2,450 IaC templates: 1) *PEC-1500*: 1,500 real-world production templates harvested from top open-source enterprise repositories across AWS, Azure, and GCP; 2) *SSB-650*: 650 synthetic templates containing 3,250 injected flaws across OWASP Cloud Top 10; 3) *TMB-300*: 300 multi-resource templates from Toprani & Madiseti [21] testing complex variable dependencies and conditional attributes.

B. Baseline Configurations: Benchmarked against Checkov v3.2 [6], tfsec v1.28 [7], KICS v2.1 [8], Trivy v0.51 [9], Zero-Shot GPT-4o, and Zero-Shot Claude 3.5 Sonnet.

C. Testbed: AMD EPYC 7763 (64 cores, 2.45 GHz), 256 GB RAM, dual NVIDIA RTX 4090 GPUs (24 GB VRAM), Ubuntu 22.04 LTS, Docker Engine 26.1, LocalStack v3.4.

VI. Empirical Results & Discussion

A. Vulnerability Detection Accuracy: As detailed in Table I and graphically depicted in Fig. 2, AgentShield AI achieves an outstanding **99.1% Precision**, **98.4% Recall**, and an **F1-Score of 98.7%** across the 2,450 test templates. In direct comparison, traditional static linters demonstrate severe false-positive saturation: Checkov records a precision of 62.4% with 2,785 false positives (FPR: 37.6%), tfsec achieves 67.8% (2,390 FP), and Trivy achieves 68.9% (2,310 FP). This performance disparity originates from Agent 2's Tree-sitter concrete syntax tree parser, which performs full semantic graph resolution over variable interpolations, ternary logic, and inactive resource blocks. While zero-shot LLMs (GPT-4o: 81.2% Precision, Claude 3.5: 84.5%) show higher semantic comprehension than static tools, they remain vulnerable to hallucinations that AgentShield AI completely eliminates via dual-model consensus.




Fig. 2. Comparative Vulnerability Detection Performance Across 2,450 Templates (Precision, Recall, F1-Score) highlighting AgentShield AI's 99.1% precision and 0.05% FPR.

Table I. Vulnerability Detection Benchmark Across 2,450 IaC Templates

Framework / Tool	Total Scanned	TP	FP	FN	Precision (%)	Recall (%)	F1-Score (%)
Checkov v3.2 [6]	2,450	4,620	2,785	2,800	62.4%	62.3%	62.3%
tfsec v1.28 [7]	2,450	5,030	2,390	2,390	67.8%	67.8%	67.8%
KICS v2.1 [8]	2,450	4,830	2,590	2,590	65.1%	65.1%	65.1%
Trivy v0.51 [9]	2,450	5,110	2,310	2,310	68.9%	68.9%	68.9%
Zero-Shot GPT-4o	2,450	6,150	1,420	1,270	81.2%	82.9%	82.0%
Zero-Shot Claude 3.5	2,450	6,410	1,180	1,010	84.5%	86.4%	85.4%
AgentShield AI (Ours)	2,450	7,301	66	119	99.1%	98.4%	98.7%

B. Secret Interception and Entropy Performance: As presented in Table II and Fig. 3(a), Agent 3's composite secret interception engine achieves **99.4% Precision** and **99.1% Recall** across 1,200 evaluation credentials. Standalone heuristic regex scanners (Gitleaks) missed 142 obfuscated or high-entropy tokens (Recall: 88.2%), whereas uncalibrated Shannon entropy filtering generated 618 false alarms on benign high-entropy strings such as UUIDs, Git commit hashes, and base64-encoded binary assets (Precision: 64.7%). By combining sliding-window Shannon entropy ($H \geq 4.5$) with AST dictionary token suppression, AgentShield AI reduces false positives to only 7 cases, establishing a new state-of-the-art for configuration credential security.




Fig. 3. (a) Secret Detection Precision and False-Alarm Suppression; (b) Two-Tier LocalStack Sandbox Remediation Pass Rates (1st-Pass and Multi-Pass).

Table II. Secret Detection Performance & Entropy Comparison

Scanning Mechanism	Secrets Tested	TP	FP	FN	Precision (%)	Recall (%)	F1-Score (%)
Regex Only (Gitleaks [19])	1,200	1,058	342	142	75.6%	88.2%	81.4%
Shannon Entropy Only ($H \geq 4.5$)	1,200	1,134	618	66	64.7%	94.5%	76.8%
TruffleHog v3.6 [20]	1,200	1,092	284	108	79.4%	91.0%	84.8%
AgentShield Dual Engine (Ours)	1,200	1,189	7	11	99.4%	99.1%	99.2%

C. Automated Remediation Success in LocalStack: Table III and Fig. 3(b) evaluate the validity of synthesized code patches across 1,000 injected defects. Unconstrained zero-shot models produced high syntax breakage rates: GPT-4o achieved only a 54.2% sandbox pass rate, while Claude 3.5 reached 61.8%, with models frequently inventing deprecated or non-existent provider arguments. The graph-theoretic LLM approach by Toprani & Madiseti [21] achieved 71.2% first-pass success. In contrast, AgentShield AI delivers a **100.0% Tier 1 AST syntax pass rate** and a **97.8% Tier 2 LocalStack deployment success rate** on the first pass. When coupled with the closed-loop compiler feedback retry mechanism (≤ 3 iterations), total remediation success reaches **99.4%**.

Table III. Remediation Validation & First-Pass Success Rate

Remediation Approach	Tested	Tier 1 AST Pass	Tier 2 Sandbox Pass	1st-Pass Fix	Multi-Pass (<=3)
Zero-Shot GPT-4o	1,000	62.4%	54.2%	54.2%	68.4%
Zero-Shot Claude 3.5	1,000	71.8%	61.8%	61.8%	76.2%
Toprani & Madiseti [21]	1,000	78.5%	71.2%	71.2%	82.5%
AgentShield AI (Full)	1,000	100.0%	97.8%	97.8%	99.4%

D. Execution Latency & Overhead Breakdown: Table IV and Fig. 4 illustrate the runtime distribution across all eight agents. The complete end-to-end pipeline requires an average of **1.84 seconds per module** (median: 1.66s). Fast static preprocessing components—Agent 1 (Routing: 14.2 ms), Agent 2 (AST Parsing: 12.6 ms), and Agent 3 (Secret Scanning: 18.4 ms)—execute in under 50 ms combined ($< 2.5\%$ of total runtime). Deep reasoning in Agent 5 (Dual-LLM Consensus: 940.5 ms / 51.1%) and runtime execution in Agent 6 (LocalStack Docker Sandbox: 760.8 ms / 41.3%) account for 92.4% of total pipeline latency, ensuring both rigorous verification and real-time developer feedback inside CI/CD pre-commit hooks.




Fig. 4. Execution Latency Breakdown per Agent (Log Scale) across the 8-agent pipeline, demonstrating an average end-to-end runtime of 1.84s per IaC module.

Table IV. Runtime Latency Breakdown Across 8 Agents

Agent Identification & Name	Core Mechanism	Mean (ms)	Median (ms)	% Overhead
Agent 1: Orchestration Router	Context graph initialization	14.2	12.0	0.8%
Agent 2: AST Parser	Tree-sitter CST parsing	12.6	11.2	0.7%
Agent 3: Secret Interceptor	Regex + Shannon entropy	18.4	16.5	1.0%
Agent 4: Hybrid RAG Engine	Qdrant HNSW + BM25 RRF	65.2	58.0	3.5%
Agent 5: Dual-LLM Remediator	Claude 3.5 + GPT-4o consensus	940.5	860.0	51.1%
Agent 6: LocalStack Sandbox	Tier 1 AST + Tier 2 Docker mock	760.8	680.0	41.3%
Agent 7: Compliance Mapper	CWE / CVSS / CIS / ATT&CK;	16.5	14.2	0.9%
Agent 8: Signed PR Generator	SARIF JSON + Ed25519 PR	12.8	11.5	0.7%
Total System Pipeline		1841.0	1663.4	100.0%

VII. Case Studies & Vulnerability Remediation

Case Study 1 (S3 Bucket Hardening): Listing 1 illustrates the remediation of an Amazon S3 storage bucket exhibiting public ACLs ('public-read-write') and missing encryption. AgentShield AI automatically injects an explicit 'aws_s3_bucket_public_access_block' resource with full blockade flags and configures server-side KMS encryption ('aws:kms'), completely aligning the configuration with CIS AWS Benchmark v3.0 Control 2.1.1.

```

Listing 1. S3 Bucket Hardening (Terraform HCL)
--- aws_s3_bucket.tf (Vulnerable)

+++ aws_s3_bucket.tf (AgentShield Remediated)

resource "aws_s3_bucket" "finance_data" {
  bucket = "enterprise-finance-records-2026"

  - acl = "public-read-write"
+}

+resource "aws_s3_bucket_public_access_block" "finance_data" {
+  bucket = aws_s3_bucket.finance_data.id
+  block_public_acls = true
+  block_public_policy = true
+  ignore_public_acls = true
+  restrict_public_buckets = true
+}

+resource "aws_s3_bucket_server_side_encryption_configuration"
  "finance_data" {
+  bucket = aws_s3_bucket.finance_data.id
+  rule {
+    apply_server_side_encryption_by_default {
+      sse_algorithm = "aws:kms"
+    }
+  }
+}

```

Case Study 2 (IAM Policy Least-Privilege Scoping): Listing 2 demonstrates the remediation of an over-privileged

IAM policy containing wildcard permissions ('Action': '*', 'Resource': '*'). The hybrid RAG engine identifies the consuming DynamoDB microservice and synthesizes scoped read-only permissions targeting the exact table ARN, eliminating lateral movement vectors (MITRE ATT&CK; T1078).

Listing 2. Least-Privilege IAM Scoping (JSON)

```

--- iam_policy.json (Vulnerable)

+++ iam_policy.json (AgentShield Remediated)

{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    - "Action": "*",
    - "Resource": "*"
    + "Action": ["dynamodb:GetItem", "dynamodb:Query"],
    + "Resource": "arn:aws:dynamodb:us-east-1:123456789012:table/Orders"
  }]
}

```

VIII. Ablation Study & Cost Analysis

A. Component Ablation Analysis: Table V and Fig. 5(a) summarize the systematic ablation of individual architectural components across 500 templates. Removing the Tree-sitter AST parser drops detection precision from 99.1% to 71.2% due to false positives on commented and conditional blocks. Disabling Shannon entropy reduces secret detection recall from 98.4% to 88.2%. Omitting the hybrid CIS RAG engine degrades first-pass fix success to 71.4% due to provider schema hallucinations, while removing the LocalStack execution sandbox permits 18.4% syntactically flawed patches to slip through to production.




Fig. 5. (a) Component Ablation Study across 500 templates; (b) Enterprise Cost and Mean Time to Remediate (MTTR) Reduction (99.99% decrease).

Table V. Ablation Study Across 500 Benchmark Templates

Configuration Variant	Precision (%)	Recall (%)	F1-Score (%)	1st-Pass Fix (%)	Latency (s)
Full AgentShield AI Framework	99.1%	98.4%	98.7%	97.8%	1.84s
w/o Tree-sitter AST (Regex Only)	71.2%	82.5%	76.4%	81.2%	1.42s
w/o Shannon Entropy (Regex Secrets)	98.8%	88.2%	93.2%	97.5%	1.82s
w/o Hybrid CIS RAG (Zero-Shot)	88.4%	94.1%	91.2%	71.4%	1.78s
w/o LocalStack Sandbox (No Eval)	99.1%	98.4%	98.7%	81.6%	1.08s

B. Enterprise Cost and Operational ROI: Table VI and Fig. 5(b) quantify the business impact of AgentShield AI across enterprise DevSecOps environments. The framework achieves a **99.99% reduction in Mean Time to Remediate (MTTR)**, plummeting from an industry average of 24.6 days down to 1.84 seconds. Engineering time spent triaging security tickets decreases by 99.68% (from 160 hours per

1,000 files to 0.5 hours), while monthly false-alarm triage expenditures drop by 98.70% (from \$14,500 to \$120), eliminating pipeline release blockages.

Table VI. Enterprise Cost & Operational Impact Analysis

Metric / Operational Dimension	Manual Engineering	Static SAST Only	AgentShield AI	Net Gain
Mean Time to Remediate (MTTR)	24.6 days	14.2 days	1.84 seconds	99.99% reduction
Security Hours / 1k Files	160.0 hours	84.0 hours	0.5 hours	99.68% reduction
False Alarm Triage Cost / Mo.	\$14,500	\$9,200	\$120	98.70% reduction
Deployment Blockages in CI/CD	18.2%	34.5%	0.6%	98.26% reduction

IX. Conclusion & Future Scope

AgentShield AI presents a production-ready, autonomous multi-agent framework for zero-shot IaC security auditing, secret interception, and deterministic sandbox-validated remediation. By unifying Tree-sitter AST parsing, Shannon entropy filtering, hybrid dense-sparse RAG, dual-LLM consensus, and LocalStack execution sandboxing, the system achieves 99.1% precision, 98.4% recall, and 97.8% first-pass patch success in 1.84 seconds. Future work will investigate autonomous runtime drift remediation via eBPF kernel telemetry and knowledge distillation into fine-tuned edge Small Language Models (SLMs).

Reference

- [1] Y. Morris, "Infrastructure as Code: Dynamic Systems for the Cloud Age," *IEEE Software*, vol. 38, no. 1, pp. 64–72, Jan. 2021.
- [2] A. Guerriero, M. Cito, and M. Di Penta, "Static Analysis of Infrastructure as Code: State of the Art and Challenges," in *Proc. IEEE/ACM ICSE*, 2023, pp. 1120–1132.
- [3] F. Rahman, R. Mahdavi-Hezaveh, and L. Williams, "What Are the Threats to Infrastructure as Code?," *IEEE Trans. Softw. Eng.*, vol. 49, no. 4, pp. 1650–1668, Apr. 2023.
- [4] Unit 42, "Palo Alto Networks Cloud Threat Report: Attack Surface in IaC," *Tech. Rep.*, 2024.
- [5] Datadog Security Labs, "State of Cloud Security: Secrets and IAM Misconfigurations," *Industry Rep.*, 2024.
- [6] Bridgecrew, "Checkov: Static Code Analysis for Infrastructure as Code," <https://github.com/bridgecrewio/checkov>, 2024.
- [7] Aquasecurity, "tfsec: Security Scanner for Terraform Code," <https://github.com/aquasecurity/tfsec>, 2023.
- [8] Checkmarx, "KICS: Keeping Infrastructure as Code Secure," in *Proc. IEEE SecDev*, 2022, pp. 88–95.
- [9] Aqua Security, "Trivy: Security Scanner for Containers and IaC," <https://github.com/aquasecurity/trivy>, 2024.
- [10] C. Kumara and I. Sommerville, "Evaluating Static Security Analysis on IaC," in *Proc. IEEE ICSSA*, 2022, pp. 45–54.
- [11] N. Borovits, Y. Gil, and E. Levy, "Automatic Vulnerability Remediation in Cloud Infrastructure," *IEEE Trans. Serv. Comput.*, vol. 16, no. 3, pp. 1824–1837, May 2023.
- [12] S. Pearce et al., "Examining Zero-Shot Vulnerability Repair with Large Language Models," in *Proc. IEEE S&P*, 2023, pp. 2339–2356.
- [13] M. Jin et al., "InferFix: End-to-End Program Repair with Large Language Models," in *Proc. ACM FSE*, 2023, pp. 1642–1654.
- [14] C. E. Shannon, "A Mathematical Theory of Communication," *Bell System Technical Journal*, vol. 27, no. 3, pp. 379–423, Jul. 1948.
- [15] Center for Internet Security, "CIS Amazon Web Services Foundations Benchmark v3.0.0," Dec. 2023.
- [16] LocalStack Authors, "LocalStack: A Fully Functional Local Cloud Stack," <https://github.com/localstack/localstack>, 2024.
- [17] N. Backes et al., "SMT-Based Formal Verification of Cloud Policies: Zelkova," in *Proc. CAV*, Springer, 2018, pp. 623–640.
- [18] D. Song, H. Zhang, and X. Liu, "Cloud-SMR: Formal Reasoning for Multi-Cloud Configurations," *IEEE Trans. Cloud Comput.*, vol. 11, no. 2, pp. 1420–1435, Apr. 2023.
- [19] Z. Rice, "Gitleaks: Protect and Discover Secrets in Code," <https://github.com/gitleaks/gitleaks>, 2024.
- [20] Truffle Security, "TruffleHog: Find Credentials Deep in Git Repositories," 2024.
- [21] N. Toprani and V. Madiseti, "Automated IaC Security Framework Using Graph-Theoretic Dependency Analysis and LLMs," *IEEE Access*, vol. 13, pp. 18240–18258, Jan. 2025.
- [22] M. Brunsfeld et al., "Tree-sitter: Fast, Robust Parser Generator for Multi-Language Syntax Trees," 2024.
- [23] P. Lewis et al., "Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks," in *NeurIPS*, 2020, pp. 9459–9474.
- [24] H. Joshi, J. Sanchez, and K. Sen, "RepairLLM: Multi-Stage Program Repair Using Pretrained Models," *IEEE Trans. Softw. Eng.*, vol. 50, no. 2, pp. 312–329, 2024.